

nizer-simple-xgb-classifier-6ed5c1

April 7, 2023

```
[1]: import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score
from xgboost import XGBClassifier

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```

```
/kaggle/input/digit-recognizer/sample_submission.csv
/kaggle/input/digit-recognizer/train.csv
/kaggle/input/digit-recognizer/test.csv
```

```
[2]: df_train = pd.read_csv("/kaggle/input/digit-recognizer/train.csv")
df_test = pd.read_csv("/kaggle/input/digit-recognizer/test.csv")
```

```
[3]: df_train
```

```
[3]:
```

	label	pixel0	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	\
0	1	0	0	0	0	0	0	0	0	
1	0	0	0	0	0	0	0	0	0	
2	1	0	0	0	0	0	0	0	0	
3	4	0	0	0	0	0	0	0	0	
4	0	0	0	0	0	0	0	0	0	
...	
41995	0	0	0	0	0	0	0	0	0	
41996	1	0	0	0	0	0	0	0	0	
41997	7	0	0	0	0	0	0	0	0	
41998	6	0	0	0	0	0	0	0	0	
41999	9	0	0	0	0	0	0	0	0	
		pixel8	...	pixel774	pixel775	pixel776	pixel777	pixel778	\	

0	0	...	0	0	0	0	0
1	0	...	0	0	0	0	0
2	0	...	0	0	0	0	0
3	0	...	0	0	0	0	0
4	0	...	0	0	0	0	0
...
41995	0	...	0	0	0	0	0
41996	0	...	0	0	0	0	0
41997	0	...	0	0	0	0	0
41998	0	...	0	0	0	0	0
41999	0	...	0	0	0	0	0

	pixel779	pixel780	pixel781	pixel782	pixel783
0	0	0	0	0	0
1	0	0	0	0	0
2	0	0	0	0	0
3	0	0	0	0	0
4	0	0	0	0	0
...
41995	0	0	0	0	0
41996	0	0	0	0	0
41997	0	0	0	0	0
41998	0	0	0	0	0
41999	0	0	0	0	0

[42000 rows x 785 columns]

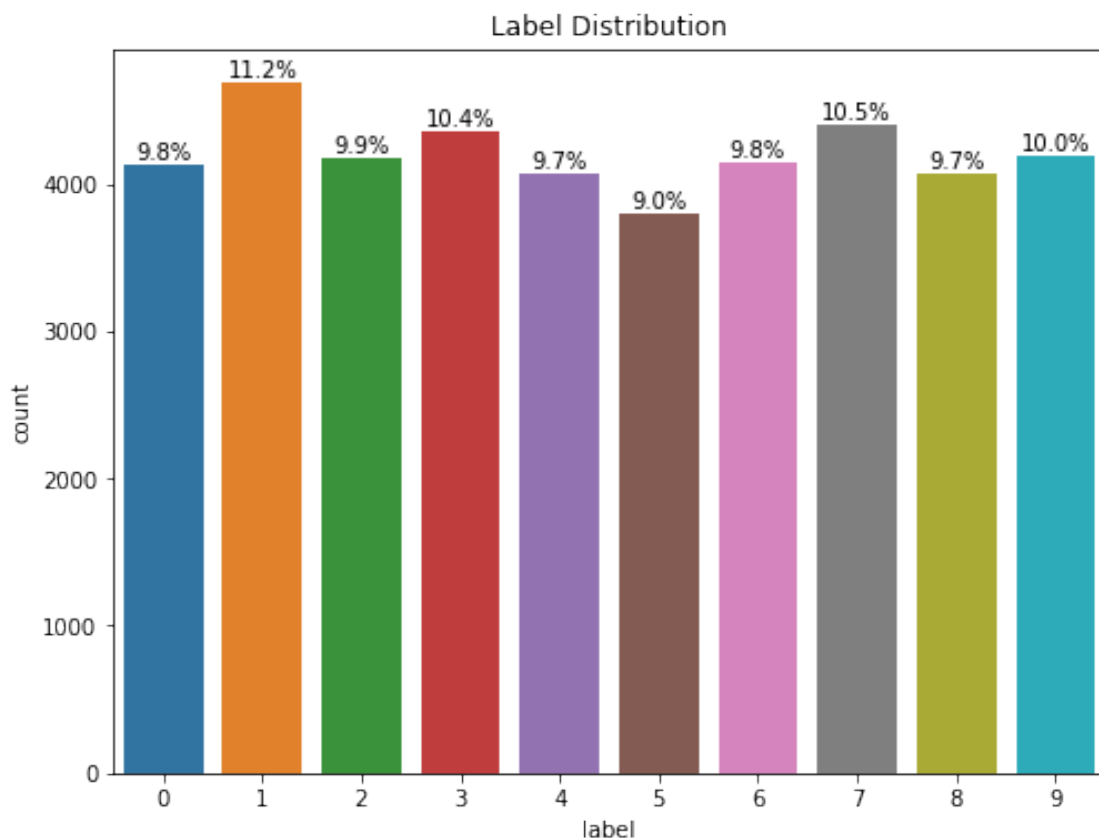
```
[4]: df_train.label.unique()
```

```
[4]: array([1, 0, 4, 7, 3, 5, 8, 9, 2, 6])
```

1 Explanatory Data Analysis

```
[5]: plt.figure(figsize=(8,6))
ax = sns.countplot(x='label',data=df_train)

plt.title("Label Distribution")
total= len(df_train.label)
for p in ax.patches:
    percentage = f'{100 * p.get_height() / total:.1f}%\n'
    x = p.get_x() + p.get_width() / 2
    y = p.get_height()
    ax.annotate(percentage, (x, y), ha='center', va='center')
```



```
[6]: df_train.describe()
```

```
[6]:
```

	label	pixel0	pixel1	pixel2	pixel3	pixel4	pixel5 \
count	42000.000000	42000.0	42000.0	42000.0	42000.0	42000.0	42000.0
mean	4.456643	0.0	0.0	0.0	0.0	0.0	0.0
std	2.887730	0.0	0.0	0.0	0.0	0.0	0.0
min	0.000000	0.0	0.0	0.0	0.0	0.0	0.0
25%	2.000000	0.0	0.0	0.0	0.0	0.0	0.0
50%	4.000000	0.0	0.0	0.0	0.0	0.0	0.0
75%	7.000000	0.0	0.0	0.0	0.0	0.0	0.0
max	9.000000	0.0	0.0	0.0	0.0	0.0	0.0

	pixel6	pixel7	pixel8 ...	pixel774	pixel775 \
count	42000.0	42000.0	42000.0 ...	42000.000000	42000.000000
mean	0.0	0.0	0.0 ...	0.219286	0.117095
std	0.0	0.0	0.0 ...	6.312890	4.633819
min	0.0	0.0	0.0 ...	0.000000	0.000000
25%	0.0	0.0	0.0 ...	0.000000	0.000000
50%	0.0	0.0	0.0 ...	0.000000	0.000000
75%	0.0	0.0	0.0 ...	0.000000	0.000000

max	0.0	0.0	0.0	...	254.000000	254.000000
-----	-----	-----	-----	-----	------------	------------

	pixel776	pixel777	pixel778	pixel779	pixel780 \
count	42000.000000	42000.000000	42000.000000	42000.000000	42000.0
mean	0.059024	0.02019	0.017238	0.002857	0.0
std	3.274488	1.75987	1.894498	0.414264	0.0
min	0.000000	0.000000	0.000000	0.000000	0.0
25%	0.000000	0.000000	0.000000	0.000000	0.0
50%	0.000000	0.000000	0.000000	0.000000	0.0
75%	0.000000	0.000000	0.000000	0.000000	0.0
max	253.000000	253.000000	254.000000	62.000000	0.0

	pixel781	pixel782	pixel783
count	42000.0	42000.0	42000.0
mean	0.0	0.0	0.0
std	0.0	0.0	0.0
min	0.0	0.0	0.0
25%	0.0	0.0	0.0
50%	0.0	0.0	0.0
75%	0.0	0.0	0.0
max	0.0	0.0	0.0

[8 rows x 785 columns]

```
[7]: df_train.sum(axis=1)
```

```
[7]: 0      16650
      1      44609
      2      13426
      3      15029
      4      51093
      ...
      41995    29310
      41996    13416
      41997    31511
      41998    26387
      41999    18187
      Length: 42000, dtype: int64
```

```
[8]: df_train.shape
```

```
[8]: (42000, 785)
```

```
[9]: pixels = df_train.columns.tolist()[1:]
      df_train["sum"] = df_train[pixels].sum(axis=1)

      df_test["sum"] = df_test[pixels].sum(axis=1)
```

```
[10]: df_train.groupby(['label'])['sum'].mean()
```

```
[10]: label
0    34632.407551
1    15188.466268
2    29871.099354
3    28320.188003
4    24232.722495
5    25835.920422
6    27734.917331
7    22931.244263
8    30184.148413
9    24553.750000
Name: sum, dtype: float64
```

```
[11]: # separate target values from df_train
targets = df_train.label
features = df_train.drop("label",axis=1)
```

```
[12]: from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
features[:] = scaler.fit_transform(features)
df_test[:] = scaler.transform(df_test)
```

```
[13]: del df_train
```

```
[14]: from sklearn.decomposition import PCA as sklearnPCA
sklearn_pca = sklearnPCA(n_components=2)
Y_sklearn = sklearn_pca.fit_transform(features)
```

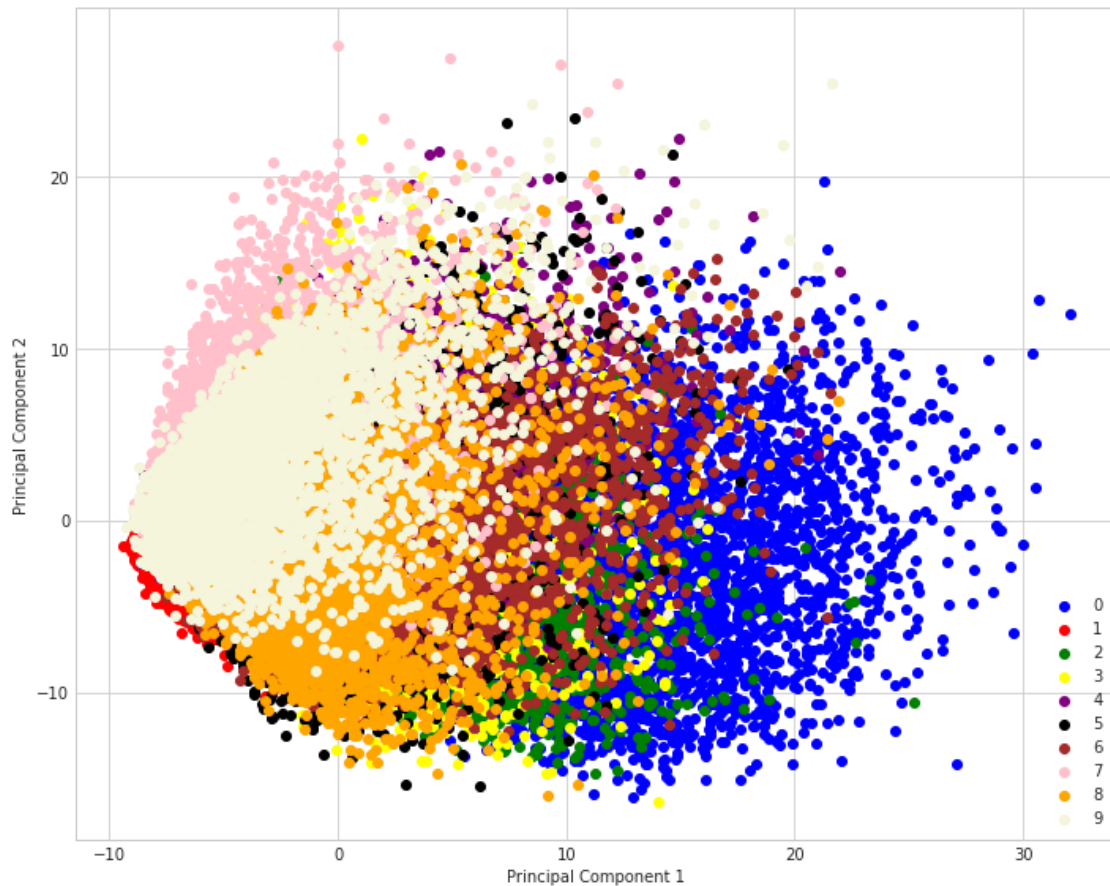
```
[15]: Y_sklearn
```

```
[15]: array([[ -5.27218318, -5.22774493],
        [19.38080907,  6.06213098],
        [-7.83436529, -1.70777172],
        ...,
        [ 0.60965064,  7.06767455],
        [ 2.25996494, -4.33643289],
        [-4.89810691,  1.55410746]])
```

```
[16]: #referred to https://sebastianraschka.com/Articles/2015\_pca\_in\_3\_steps.html and
      ↪ https://www.kaggle.com/arthurtok/
      ↪ interactive-intro-to-dimensionality-reduction

with plt.style.context('seaborn-whitegrid'):
```

```
plt.figure(figsize=(10, 8))
for lab, col in zip((0,1,2,3,4,5,6,7,8,9),
                    ('blue', 'red', 'green', 'yellow', 'purple', 'black', 'brown', 'pink', 'orange', 'beige')):
    plt.scatter(Y_sklern[target==lab, 0],
                Y_sklern[target==lab, 1],
                label=lab,
                c=col)
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.legend(loc='lower right')
plt.tight_layout()
plt.show()
```



```
[17]: features.index
```

```
[17]: RangeIndex(start=0, stop=42000, step=1)
```

```
[18]: sklearn_pca_3 = sklearnPCA(n_components=3)
Y_sklearn_3 = sklearn_pca_3.fit_transform(features)
Y_sklearn_3_test = sklearn_pca_3.transform(df_test)
```

```
[19]: # Store results of PCA in a data frame
result=pd.DataFrame(Y_sklearn_3, columns=['PCA%i' % i for i in range(3)],
↳ index=features.index)
```

```
[20]: result
```

```
[20]:
```

	PCA0	PCA1	PCA2
0	-5.272200	-5.227576	3.888980
1	19.380803	6.062140	1.339064
2	-7.834355	-1.707918	2.291291
3	-0.706307	5.845857	2.023411
4	26.648676	6.067638	0.981800
...
41995	13.527937	-1.320166	-3.914860
41996	-9.041423	-1.192330	2.320770
41997	0.609672	7.067932	-12.100110
41998	2.259955	-4.337248	0.714577
41999	-4.898145	1.554266	-2.501381

[42000 rows x 3 columns]

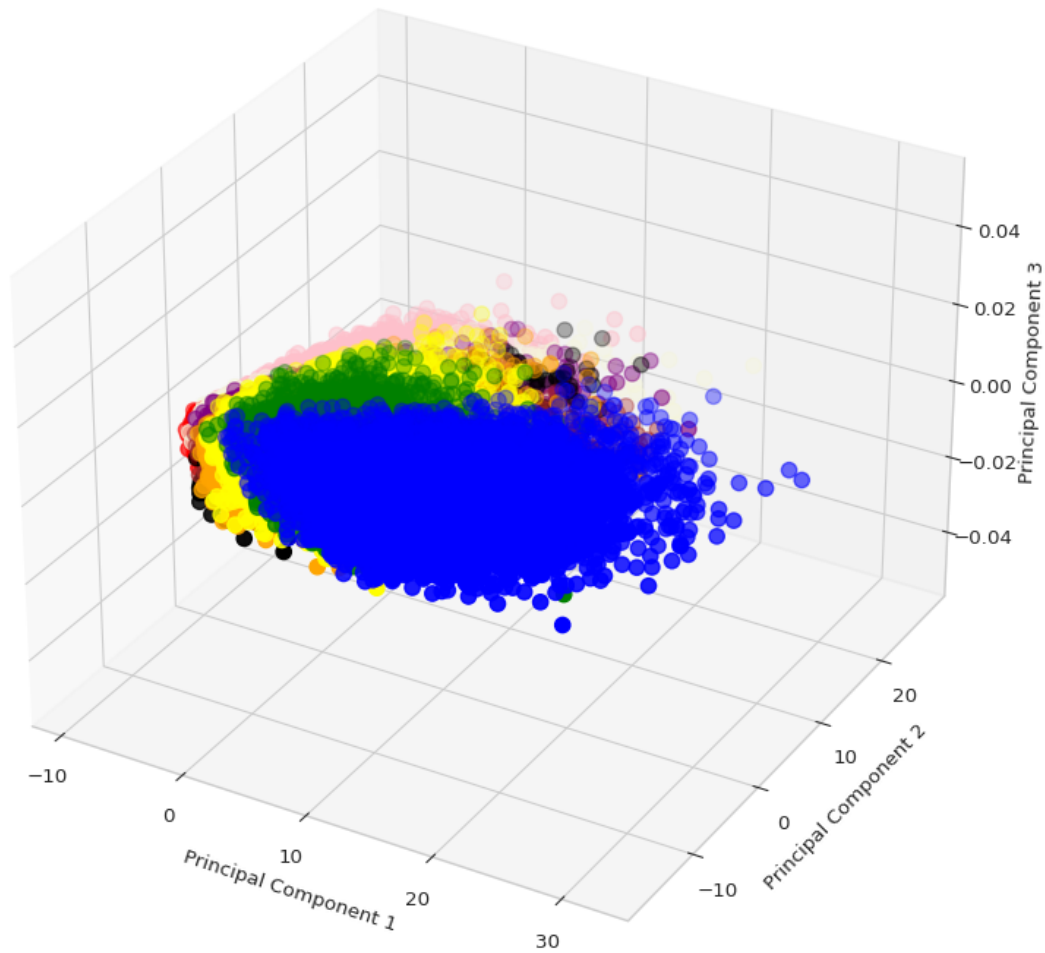
```
[21]: my_dpi=96
plt.figure(figsize=(480/my_dpi, 480/my_dpi), dpi=my_dpi)

with plt.style.context('seaborn-whitegrid'):
    my_dpi=96
    fig = plt.figure(figsize=(10, 10), dpi=my_dpi)
    ax = fig.add_subplot(111,projection = '3d')
    for lab, col in zip((0,1,2,3,4,5,6,7,8,9),
↳ ('blue', 'red', 'green', 'yellow', 'purple', 'black', 'brown', 'pink', 'orange', 'beige')):
↳
        plt.scatter(Y_sklearn[target==lab, 0],
                    Y_sklearn[target==lab, 1],
                    label=lab,
                    c=col,s =60)

    ax.set_xlabel('Principal Component 1')
    ax.set_ylabel('Principal Component 2')
    ax.set_zlabel('Principal Component 3')
    ax.set_title("PCA on the Handwriting Data")
    plt.show()
```

<Figure size 480x480 with 0 Axes>

PCA on the Handwriting Data



```
[22]: encoder = LabelEncoder()  
      targets[:] = encoder.fit_transform(targets[:])
```

```
[23]: X_train,X_val, y_train,y_val = train_test_split(result,targets,random_state=1)
```

2 Making a Model and Predictions

```
[24]: # 3 Principal Components  
      model = XGBClassifier(max_depth=5, objective='multi:softprob',  
                           ↪n_estimators=1000,
```



```

num_classes=10)

history = model.fit(X_train, y_train, eval_set=
    ↪=[(X_val, y_val)], early_stopping_rounds=50)
acc = accuracy_score(y_val, model.predict(X_val))
print(f"Accuracy: , {round(acc,3)}")

```

/opt/conda/lib/python3.7/site-packages/xgboost/sklearn.py:1224: UserWarning: The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option use_label_encoder=False when constructing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].
warnings.warn(label_encoder_deprecation_msg, UserWarning)

[11:37:36] WARNING: ../src/learner.cc:576:
Parameters: { "num_classes" } might not be used.

This could be a false alarm, with some parameters getting used by language bindings but
then being mistakenly passed down to XGBoost core, or some parameter actually being used
but getting flagged wrongly here. Please open an issue if you find any such cases.

[11:37:36] WARNING: ../src/learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

```

[0]    validation_0-mlogloss:1.87620
[1]    validation_0-mlogloss:1.69532
[2]    validation_0-mlogloss:1.57450
[3]    validation_0-mlogloss:1.49041
[4]    validation_0-mlogloss:1.42499
[5]    validation_0-mlogloss:1.37516
[6]    validation_0-mlogloss:1.33594
[7]    validation_0-mlogloss:1.30528
[8]    validation_0-mlogloss:1.28138
[9]    validation_0-mlogloss:1.26042
[10]   validation_0-mlogloss:1.24410
[11]   validation_0-mlogloss:1.23011
[12]   validation_0-mlogloss:1.21839
[13]   validation_0-mlogloss:1.20846
[14]   validation_0-mlogloss:1.20028
[15]   validation_0-mlogloss:1.19254
[16]   validation_0-mlogloss:1.18694
[17]   validation_0-mlogloss:1.18069
[18]   validation_0-mlogloss:1.17630

```

[19] validation_0-mlogloss:1.17323
[20] validation_0-mlogloss:1.17057
[21] validation_0-mlogloss:1.16672
[22] validation_0-mlogloss:1.16426
[23] validation_0-mlogloss:1.16236
[24] validation_0-mlogloss:1.16045
[25] validation_0-mlogloss:1.15898
[26] validation_0-mlogloss:1.15801
[27] validation_0-mlogloss:1.15613
[28] validation_0-mlogloss:1.15503
[29] validation_0-mlogloss:1.15392
[30] validation_0-mlogloss:1.15284
[31] validation_0-mlogloss:1.15198
[32] validation_0-mlogloss:1.15184
[33] validation_0-mlogloss:1.15075
[34] validation_0-mlogloss:1.15035
[35] validation_0-mlogloss:1.14970
[36] validation_0-mlogloss:1.14949
[37] validation_0-mlogloss:1.14931
[38] validation_0-mlogloss:1.14892
[39] validation_0-mlogloss:1.14844
[40] validation_0-mlogloss:1.14843
[41] validation_0-mlogloss:1.14848
[42] validation_0-mlogloss:1.14835
[43] validation_0-mlogloss:1.14812
[44] validation_0-mlogloss:1.14821
[45] validation_0-mlogloss:1.14781
[46] validation_0-mlogloss:1.14776
[47] validation_0-mlogloss:1.14804
[48] validation_0-mlogloss:1.14821
[49] validation_0-mlogloss:1.14820
[50] validation_0-mlogloss:1.14821
[51] validation_0-mlogloss:1.14821
[52] validation_0-mlogloss:1.14837
[53] validation_0-mlogloss:1.14841
[54] validation_0-mlogloss:1.14838
[55] validation_0-mlogloss:1.14846
[56] validation_0-mlogloss:1.14864
[57] validation_0-mlogloss:1.14880
[58] validation_0-mlogloss:1.14898
[59] validation_0-mlogloss:1.14932
[60] validation_0-mlogloss:1.14910
[61] validation_0-mlogloss:1.14920
[62] validation_0-mlogloss:1.14941
[63] validation_0-mlogloss:1.14962
[64] validation_0-mlogloss:1.14978
[65] validation_0-mlogloss:1.14941
[66] validation_0-mlogloss:1.14915

```

[67] validation_0-mlogloss:1.14919
[68] validation_0-mlogloss:1.14951
[69] validation_0-mlogloss:1.14995
[70] validation_0-mlogloss:1.15012
[71] validation_0-mlogloss:1.15042
[72] validation_0-mlogloss:1.15072
[73] validation_0-mlogloss:1.15102
[74] validation_0-mlogloss:1.15121
[75] validation_0-mlogloss:1.15103
[76] validation_0-mlogloss:1.15137
[77] validation_0-mlogloss:1.15157
[78] validation_0-mlogloss:1.15196
[79] validation_0-mlogloss:1.15205
[80] validation_0-mlogloss:1.15232
[81] validation_0-mlogloss:1.15233
[82] validation_0-mlogloss:1.15241
[83] validation_0-mlogloss:1.15252
[84] validation_0-mlogloss:1.15278
[85] validation_0-mlogloss:1.15281
[86] validation_0-mlogloss:1.15290
[87] validation_0-mlogloss:1.15296
[88] validation_0-mlogloss:1.15307
[89] validation_0-mlogloss:1.15325
[90] validation_0-mlogloss:1.15331
[91] validation_0-mlogloss:1.15385
[92] validation_0-mlogloss:1.15409
[93] validation_0-mlogloss:1.15446
[94] validation_0-mlogloss:1.15469
[95] validation_0-mlogloss:1.15475

```

Accuracy: , 0.559

```
[25]: X_train,X_val, y_train,y_val = train_test_split(features,targets,random_state=1)
```

```
[26]: model = XGBClassifier(max_depth=5, objective='multi:softprob',
    ↪n_estimators=1000,
    num_classes=10)

history = model.fit(X_train, y_train,eval_set=
    ↪[(X_train,y_train),(X_val,y_val)],early_stopping_rounds =5)
acc = accuracy_score(y_val, model.predict(X_val))
print(f"Accuracy: , {round(acc,3)}")
```

```
[11:37:52] WARNING: ../src/learner.cc:576:
Parameters: { "num_classes" } might not be used.
```

This could be a false alarm, with some parameters getting used by language bindings but then being mistakenly passed down to XGBoost core, or some parameter actually

being used

but getting flagged wrongly here. Please open an issue if you find any such cases.

[11:37:56] WARNING: ../src/learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[0]	validation_0-mlogloss:1.42839	validation_1-mlogloss:1.44561
[1]	validation_0-mlogloss:1.09931	validation_1-mlogloss:1.13044
[2]	validation_0-mlogloss:0.88151	validation_1-mlogloss:0.92129
[3]	validation_0-mlogloss:0.72542	validation_1-mlogloss:0.77069
[4]	validation_0-mlogloss:0.60957	validation_1-mlogloss:0.65913
[5]	validation_0-mlogloss:0.51863	validation_1-mlogloss:0.57078
[6]	validation_0-mlogloss:0.44818	validation_1-mlogloss:0.50382
[7]	validation_0-mlogloss:0.38978	validation_1-mlogloss:0.44809
[8]	validation_0-mlogloss:0.34231	validation_1-mlogloss:0.40327
[9]	validation_0-mlogloss:0.30416	validation_1-mlogloss:0.36713
[10]	validation_0-mlogloss:0.27249	validation_1-mlogloss:0.33768
[11]	validation_0-mlogloss:0.24479	validation_1-mlogloss:0.31212
[12]	validation_0-mlogloss:0.22142	validation_1-mlogloss:0.29098
[13]	validation_0-mlogloss:0.20190	validation_1-mlogloss:0.27290
[14]	validation_0-mlogloss:0.18458	validation_1-mlogloss:0.25708
[15]	validation_0-mlogloss:0.16888	validation_1-mlogloss:0.24343
[16]	validation_0-mlogloss:0.15593	validation_1-mlogloss:0.23174
[17]	validation_0-mlogloss:0.14441	validation_1-mlogloss:0.22166
[18]	validation_0-mlogloss:0.13361	validation_1-mlogloss:0.21180
[19]	validation_0-mlogloss:0.12413	validation_1-mlogloss:0.20371
[20]	validation_0-mlogloss:0.11625	validation_1-mlogloss:0.19627
[21]	validation_0-mlogloss:0.10815	validation_1-mlogloss:0.18900
[22]	validation_0-mlogloss:0.10031	validation_1-mlogloss:0.18205
[23]	validation_0-mlogloss:0.09473	validation_1-mlogloss:0.17735
[24]	validation_0-mlogloss:0.08789	validation_1-mlogloss:0.17083
[25]	validation_0-mlogloss:0.08249	validation_1-mlogloss:0.16605
[26]	validation_0-mlogloss:0.07738	validation_1-mlogloss:0.16129
[27]	validation_0-mlogloss:0.07189	validation_1-mlogloss:0.15697
[28]	validation_0-mlogloss:0.06801	validation_1-mlogloss:0.15350
[29]	validation_0-mlogloss:0.06414	validation_1-mlogloss:0.14999
[30]	validation_0-mlogloss:0.06057	validation_1-mlogloss:0.14722
[31]	validation_0-mlogloss:0.05713	validation_1-mlogloss:0.14373
[32]	validation_0-mlogloss:0.05349	validation_1-mlogloss:0.14027
[33]	validation_0-mlogloss:0.05094	validation_1-mlogloss:0.13833
[34]	validation_0-mlogloss:0.04831	validation_1-mlogloss:0.13593
[35]	validation_0-mlogloss:0.04560	validation_1-mlogloss:0.13382
[36]	validation_0-mlogloss:0.04306	validation_1-mlogloss:0.13161
[37]	validation_0-mlogloss:0.04033	validation_1-mlogloss:0.12903
[38]	validation_0-mlogloss:0.03796	validation_1-mlogloss:0.12721

[39]	validation_0-mlogloss:0.03598	validation_1-mlogloss:0.12509
[40]	validation_0-mlogloss:0.03401	validation_1-mlogloss:0.12340
[41]	validation_0-mlogloss:0.03242	validation_1-mlogloss:0.12218
[42]	validation_0-mlogloss:0.03080	validation_1-mlogloss:0.12055
[43]	validation_0-mlogloss:0.02905	validation_1-mlogloss:0.11884
[44]	validation_0-mlogloss:0.02748	validation_1-mlogloss:0.11706
[45]	validation_0-mlogloss:0.02617	validation_1-mlogloss:0.11563
[46]	validation_0-mlogloss:0.02485	validation_1-mlogloss:0.11447
[47]	validation_0-mlogloss:0.02354	validation_1-mlogloss:0.11336
[48]	validation_0-mlogloss:0.02231	validation_1-mlogloss:0.11205
[49]	validation_0-mlogloss:0.02120	validation_1-mlogloss:0.11124
[50]	validation_0-mlogloss:0.02023	validation_1-mlogloss:0.10998
[51]	validation_0-mlogloss:0.01928	validation_1-mlogloss:0.10883
[52]	validation_0-mlogloss:0.01829	validation_1-mlogloss:0.10787
[53]	validation_0-mlogloss:0.01723	validation_1-mlogloss:0.10640
[54]	validation_0-mlogloss:0.01643	validation_1-mlogloss:0.10563
[55]	validation_0-mlogloss:0.01576	validation_1-mlogloss:0.10474
[56]	validation_0-mlogloss:0.01502	validation_1-mlogloss:0.10384
[57]	validation_0-mlogloss:0.01445	validation_1-mlogloss:0.10332
[58]	validation_0-mlogloss:0.01378	validation_1-mlogloss:0.10245
[59]	validation_0-mlogloss:0.01308	validation_1-mlogloss:0.10127
[60]	validation_0-mlogloss:0.01244	validation_1-mlogloss:0.10036
[61]	validation_0-mlogloss:0.01200	validation_1-mlogloss:0.09975
[62]	validation_0-mlogloss:0.01149	validation_1-mlogloss:0.09906
[63]	validation_0-mlogloss:0.01097	validation_1-mlogloss:0.09831
[64]	validation_0-mlogloss:0.01044	validation_1-mlogloss:0.09763
[65]	validation_0-mlogloss:0.01000	validation_1-mlogloss:0.09712
[66]	validation_0-mlogloss:0.00961	validation_1-mlogloss:0.09666
[67]	validation_0-mlogloss:0.00912	validation_1-mlogloss:0.09616
[68]	validation_0-mlogloss:0.00872	validation_1-mlogloss:0.09550
[69]	validation_0-mlogloss:0.00829	validation_1-mlogloss:0.09468
[70]	validation_0-mlogloss:0.00795	validation_1-mlogloss:0.09430
[71]	validation_0-mlogloss:0.00766	validation_1-mlogloss:0.09384
[72]	validation_0-mlogloss:0.00734	validation_1-mlogloss:0.09339
[73]	validation_0-mlogloss:0.00706	validation_1-mlogloss:0.09305
[74]	validation_0-mlogloss:0.00676	validation_1-mlogloss:0.09246
[75]	validation_0-mlogloss:0.00646	validation_1-mlogloss:0.09196
[76]	validation_0-mlogloss:0.00620	validation_1-mlogloss:0.09154
[77]	validation_0-mlogloss:0.00600	validation_1-mlogloss:0.09107
[78]	validation_0-mlogloss:0.00578	validation_1-mlogloss:0.09086
[79]	validation_0-mlogloss:0.00557	validation_1-mlogloss:0.09031
[80]	validation_0-mlogloss:0.00537	validation_1-mlogloss:0.09002
[81]	validation_0-mlogloss:0.00517	validation_1-mlogloss:0.08971
[82]	validation_0-mlogloss:0.00500	validation_1-mlogloss:0.08952
[83]	validation_0-mlogloss:0.00484	validation_1-mlogloss:0.08928
[84]	validation_0-mlogloss:0.00464	validation_1-mlogloss:0.08888
[85]	validation_0-mlogloss:0.00451	validation_1-mlogloss:0.08865
[86]	validation_0-mlogloss:0.00435	validation_1-mlogloss:0.08854

[87]	validation_0-mlogloss:0.00424	validation_1-mlogloss:0.08818
[88]	validation_0-mlogloss:0.00410	validation_1-mlogloss:0.08801
[89]	validation_0-mlogloss:0.00398	validation_1-mlogloss:0.08777
[90]	validation_0-mlogloss:0.00383	validation_1-mlogloss:0.08748
[91]	validation_0-mlogloss:0.00371	validation_1-mlogloss:0.08738
[92]	validation_0-mlogloss:0.00359	validation_1-mlogloss:0.08702
[93]	validation_0-mlogloss:0.00348	validation_1-mlogloss:0.08697
[94]	validation_0-mlogloss:0.00338	validation_1-mlogloss:0.08668
[95]	validation_0-mlogloss:0.00327	validation_1-mlogloss:0.08649
[96]	validation_0-mlogloss:0.00318	validation_1-mlogloss:0.08631
[97]	validation_0-mlogloss:0.00309	validation_1-mlogloss:0.08618
[98]	validation_0-mlogloss:0.00301	validation_1-mlogloss:0.08595
[99]	validation_0-mlogloss:0.00294	validation_1-mlogloss:0.08581
[100]	validation_0-mlogloss:0.00287	validation_1-mlogloss:0.08575
[101]	validation_0-mlogloss:0.00279	validation_1-mlogloss:0.08545
[102]	validation_0-mlogloss:0.00271	validation_1-mlogloss:0.08522
[103]	validation_0-mlogloss:0.00264	validation_1-mlogloss:0.08503
[104]	validation_0-mlogloss:0.00258	validation_1-mlogloss:0.08492
[105]	validation_0-mlogloss:0.00252	validation_1-mlogloss:0.08472
[106]	validation_0-mlogloss:0.00246	validation_1-mlogloss:0.08459
[107]	validation_0-mlogloss:0.00240	validation_1-mlogloss:0.08452
[108]	validation_0-mlogloss:0.00235	validation_1-mlogloss:0.08442
[109]	validation_0-mlogloss:0.00230	validation_1-mlogloss:0.08439
[110]	validation_0-mlogloss:0.00225	validation_1-mlogloss:0.08441
[111]	validation_0-mlogloss:0.00220	validation_1-mlogloss:0.08430
[112]	validation_0-mlogloss:0.00215	validation_1-mlogloss:0.08411
[113]	validation_0-mlogloss:0.00210	validation_1-mlogloss:0.08391
[114]	validation_0-mlogloss:0.00205	validation_1-mlogloss:0.08373
[115]	validation_0-mlogloss:0.00202	validation_1-mlogloss:0.08373
[116]	validation_0-mlogloss:0.00197	validation_1-mlogloss:0.08364
[117]	validation_0-mlogloss:0.00194	validation_1-mlogloss:0.08346
[118]	validation_0-mlogloss:0.00190	validation_1-mlogloss:0.08344
[119]	validation_0-mlogloss:0.00187	validation_1-mlogloss:0.08342
[120]	validation_0-mlogloss:0.00183	validation_1-mlogloss:0.08336
[121]	validation_0-mlogloss:0.00179	validation_1-mlogloss:0.08324
[122]	validation_0-mlogloss:0.00176	validation_1-mlogloss:0.08306
[123]	validation_0-mlogloss:0.00172	validation_1-mlogloss:0.08294
[124]	validation_0-mlogloss:0.00169	validation_1-mlogloss:0.08292
[125]	validation_0-mlogloss:0.00166	validation_1-mlogloss:0.08289
[126]	validation_0-mlogloss:0.00164	validation_1-mlogloss:0.08286
[127]	validation_0-mlogloss:0.00161	validation_1-mlogloss:0.08281
[128]	validation_0-mlogloss:0.00158	validation_1-mlogloss:0.08274
[129]	validation_0-mlogloss:0.00155	validation_1-mlogloss:0.08273
[130]	validation_0-mlogloss:0.00153	validation_1-mlogloss:0.08260
[131]	validation_0-mlogloss:0.00150	validation_1-mlogloss:0.08258
[132]	validation_0-mlogloss:0.00148	validation_1-mlogloss:0.08249
[133]	validation_0-mlogloss:0.00145	validation_1-mlogloss:0.08250
[134]	validation_0-mlogloss:0.00143	validation_1-mlogloss:0.08239

[135]	validation_0-mlogloss:0.00141	validation_1-mlogloss:0.08235
[136]	validation_0-mlogloss:0.00138	validation_1-mlogloss:0.08235
[137]	validation_0-mlogloss:0.00136	validation_1-mlogloss:0.08233
[138]	validation_0-mlogloss:0.00134	validation_1-mlogloss:0.08227
[139]	validation_0-mlogloss:0.00132	validation_1-mlogloss:0.08229
[140]	validation_0-mlogloss:0.00130	validation_1-mlogloss:0.08223
[141]	validation_0-mlogloss:0.00129	validation_1-mlogloss:0.08227
[142]	validation_0-mlogloss:0.00127	validation_1-mlogloss:0.08219
[143]	validation_0-mlogloss:0.00125	validation_1-mlogloss:0.08212
[144]	validation_0-mlogloss:0.00123	validation_1-mlogloss:0.08204
[145]	validation_0-mlogloss:0.00122	validation_1-mlogloss:0.08200
[146]	validation_0-mlogloss:0.00120	validation_1-mlogloss:0.08204
[147]	validation_0-mlogloss:0.00119	validation_1-mlogloss:0.08193
[148]	validation_0-mlogloss:0.00117	validation_1-mlogloss:0.08193
[149]	validation_0-mlogloss:0.00116	validation_1-mlogloss:0.08186
[150]	validation_0-mlogloss:0.00114	validation_1-mlogloss:0.08187
[151]	validation_0-mlogloss:0.00113	validation_1-mlogloss:0.08187
[152]	validation_0-mlogloss:0.00111	validation_1-mlogloss:0.08185
[153]	validation_0-mlogloss:0.00110	validation_1-mlogloss:0.08190
[154]	validation_0-mlogloss:0.00109	validation_1-mlogloss:0.08171
[155]	validation_0-mlogloss:0.00107	validation_1-mlogloss:0.08170
[156]	validation_0-mlogloss:0.00106	validation_1-mlogloss:0.08162
[157]	validation_0-mlogloss:0.00105	validation_1-mlogloss:0.08158
[158]	validation_0-mlogloss:0.00104	validation_1-mlogloss:0.08154
[159]	validation_0-mlogloss:0.00103	validation_1-mlogloss:0.08149
[160]	validation_0-mlogloss:0.00101	validation_1-mlogloss:0.08153
[161]	validation_0-mlogloss:0.00100	validation_1-mlogloss:0.08147
[162]	validation_0-mlogloss:0.00099	validation_1-mlogloss:0.08142
[163]	validation_0-mlogloss:0.00098	validation_1-mlogloss:0.08138
[164]	validation_0-mlogloss:0.00097	validation_1-mlogloss:0.08144
[165]	validation_0-mlogloss:0.00096	validation_1-mlogloss:0.08140
[166]	validation_0-mlogloss:0.00095	validation_1-mlogloss:0.08142
[167]	validation_0-mlogloss:0.00094	validation_1-mlogloss:0.08137
[168]	validation_0-mlogloss:0.00093	validation_1-mlogloss:0.08130
[169]	validation_0-mlogloss:0.00092	validation_1-mlogloss:0.08129
[170]	validation_0-mlogloss:0.00091	validation_1-mlogloss:0.08131
[171]	validation_0-mlogloss:0.00090	validation_1-mlogloss:0.08130
[172]	validation_0-mlogloss:0.00090	validation_1-mlogloss:0.08127
[173]	validation_0-mlogloss:0.00089	validation_1-mlogloss:0.08126
[174]	validation_0-mlogloss:0.00088	validation_1-mlogloss:0.08131
[175]	validation_0-mlogloss:0.00087	validation_1-mlogloss:0.08117
[176]	validation_0-mlogloss:0.00086	validation_1-mlogloss:0.08117
[177]	validation_0-mlogloss:0.00086	validation_1-mlogloss:0.08115
[178]	validation_0-mlogloss:0.00085	validation_1-mlogloss:0.08114
[179]	validation_0-mlogloss:0.00084	validation_1-mlogloss:0.08113
[180]	validation_0-mlogloss:0.00084	validation_1-mlogloss:0.08112
[181]	validation_0-mlogloss:0.00083	validation_1-mlogloss:0.08109
[182]	validation_0-mlogloss:0.00082	validation_1-mlogloss:0.08103

```

[183] validation_0-mlogloss:0.00081 validation_1-mlogloss:0.08103
[184] validation_0-mlogloss:0.00081 validation_1-mlogloss:0.08102
[185] validation_0-mlogloss:0.00080 validation_1-mlogloss:0.08100
[186] validation_0-mlogloss:0.00080 validation_1-mlogloss:0.08102
[187] validation_0-mlogloss:0.00079 validation_1-mlogloss:0.08093
[188] validation_0-mlogloss:0.00078 validation_1-mlogloss:0.08083
[189] validation_0-mlogloss:0.00078 validation_1-mlogloss:0.08081
[190] validation_0-mlogloss:0.00077 validation_1-mlogloss:0.08086
[191] validation_0-mlogloss:0.00076 validation_1-mlogloss:0.08087
[192] validation_0-mlogloss:0.00076 validation_1-mlogloss:0.08091
[193] validation_0-mlogloss:0.00075 validation_1-mlogloss:0.08097
[194] validation_0-mlogloss:0.00075 validation_1-mlogloss:0.08099
Accuracy: , 0.976

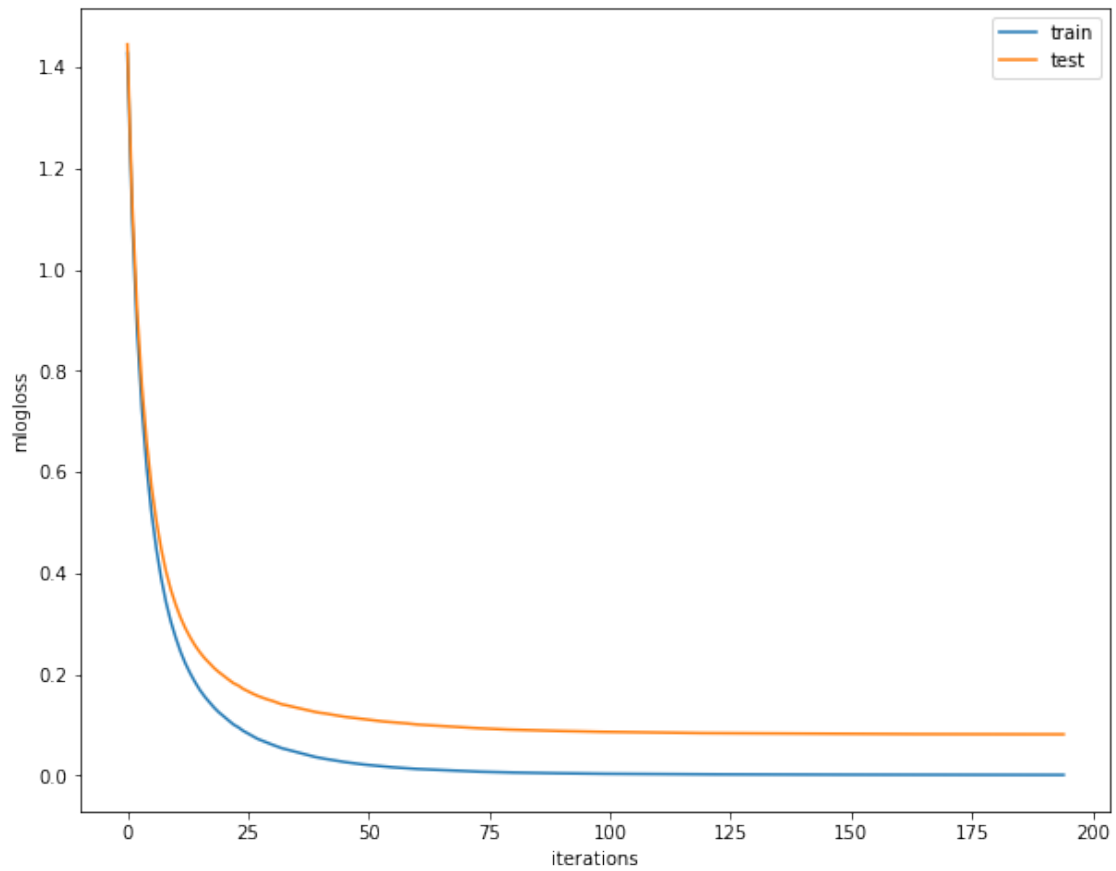
```

```
[27]: results = model.eval_result()
```

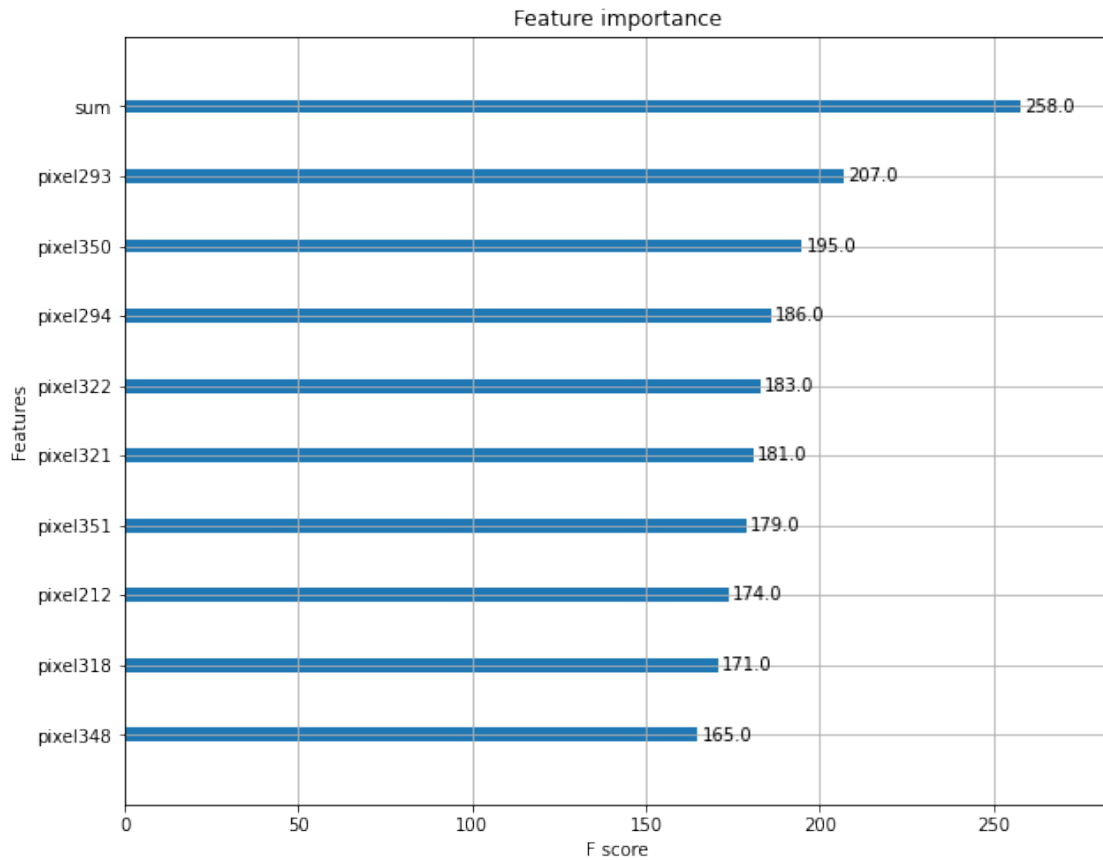
```

[28]: from matplotlib import pyplot
      # plot learning curves
      plt.figure(figsize=(10, 8))
      pyplot.plot(results['validation_0']['mlogloss'], label='train')
      pyplot.plot(results['validation_1']['mlogloss'], label='test')
      # show the legend
      pyplot.legend()
      plt.xlabel('iterations')
      plt.ylabel('mlogloss')
      # show the plot
      pyplot.show()

```

```
[29]: from xgboost import plot_importance
ax = plot_importance(model,max_num_features=10)
fig = ax.figure
fig.set_size_inches(10,8)
plt.show()
```



```
[30]: predictions = model.predict(df_test)
```

```
[31]: output = pd.read_csv("../input/digit-recognizer/sample_submission.csv")
output['Label'] = predictions
output.to_csv('submission.csv', index=False)
```

##After Learning the above Notebook one has to take print of it and answer following questions in writing and collectively complete write up.

#Q.1 What is Decision Tree Algorithm ? Which type of ML we can solve using Decision Tree?

#ANSWER:

Decision Tree is a Supervised learning technique that can be used for both classification and Regression problems, but mostly it is preferred for solving Classification problems. It is a tree-structured classifier, where internal nodes represent the features of a dataset, branches represent the decision rules and each leaf node represents the outcome. In a Decision tree, there are two nodes, which are the Decision Node and Leaf Node. Decision nodes are used to make any decision and have multiple branches, whereas Leaf nodes are the output of those decisions and do not contain any further branches. The decisions or the test are performed on the basis of features of the given dataset. It is a graphical representation for getting all the possible solutions to a problem/decision based on

given conditions. It is called a decision tree because, similar to a tree, it starts with the root node, which expands on further branches and constructs a tree-like structure.

Decision trees are a popular supervised learning method that like many other learning methods we've seen, can be used for both regression and classification. Decision trees in machine learning (ML) are used to structure algorithms. A decision tree algorithm helps split dataset features with a cost function. Through a process called pruning, the trees are grown before being optimized to remove branches that use irrelevant features.

#Q.2 What do you mean by ensemble learning ? Does XGBoost support ensemble learning ?

#ANSWER:

Ensemble learning is a general meta approach to machine learning that seeks better predictive performance by combining the predictions from multiple models. Ensemble learning is the process by which multiple models, such as classifiers or experts, are strategically generated and combined to solve a particular computational intelligence problem. Ensemble learning is primarily used to improve the (classification, prediction, function approximation, etc. There are two main reasons to use an ensemble over a single model, and they are related; they are: Performance: An ensemble can make better predictions and achieve better performance than any single contributing model. Robustness: An ensemble reduces the spread or dispersion of the predictions and model performance.

XGBoost, that is, Extreme Gradient Boosting, is a very popular machine learning ensemble technique that has helped data scientists across the globe to achieve great results with phenomenal accuracy. XGBoost is built on the principles of ensemble modeling and is an improved version of the Gradient Boosted Machine algorithm. In general, the XgBoost algorithm creates multiple classifiers that are weak learners, which means a model that gives a bit better accuracy than just a random guess. The learner in the ensemble model can be a linear or tree model that is built iteratively with random sampling along with an added weight from the learnings of the previously built model. At each step, a tree is built and the cases where the tree has failed to classify an outcome correctly is assigned a corresponding weight. The next iteration of model building learns from the mistakes of the previous model. Gradient boosting refers to a class of ensemble machine learning algorithms that can be used for classification or regression predictive modeling problems.

Ensembles are constructed from decision tree models. Trees are added one at a time to the ensemble and fit to correct the prediction errors made by prior models. This is a type of ensemble machine learning model referred to as boosting.

#Q.3 What is Principal Component Analysis ? Why do we use PCA in our notebook?

#ANSWER:

PCA is a widely covered machine learning method on the web, and there are some great articles about it, but many spend too much time in the weeds on the topic, when most of us just want to know how it works in a simplified way.

Principal component analysis can be broken down into five steps. I'll go through each step, providing logical explanations of what PCA is doing and simplifying mathematical concepts such as standardization, covariance, eigenvectors and eigenvalues without focusing on how to compute them.

PCA helps you interpret your data, but it will not always find the important patterns. Principal

component analysis (PCA) simplifies the complexity in high-dimensional data while retaining trends and patterns. It does this by transforming the data into fewer dimensions, which act as summaries of features. PCA is used to visualize multidimensional data. It is used to reduce the number of dimensions in healthcare data. PCA can help resize an image. When PCA is used as part of preprocessing, the algorithm is applied to: Reduce the number of dimensions in the training dataset. De-noise the data. Because PCA is computed by finding the components which explain the greatest amount of variance, it captures the signal in the data and omits the noise.

#Q.4 Check use of “StandardScaler” class from sklearn in notebook. What do you think is this API used for?

#ANSWER:

Python sklearn library offers us with StandardScaler() function to standardize the data values into a standard format. According to the above syntax, we initially create an object of the StandardScaler() function. Further, we use fit_transform() along with the assigned object to transform the data and standardize it. StandardScaler removes the mean and scales each feature/variable to unit variance. This operation is performed feature-wise in an independent way. StandardScaler can be influenced by outliers (if they exist in the dataset) since it involves the estimation of the empirical mean and standard deviation of each feature. When the features of the given dataset fluctuate significantly within their ranges or are recorded in various units of measurement, StandardScaler enters the picture.

The data are scaled to a variance of 1 after the mean is reduced to 0 via StandardScaler. But when determining the empirical mean of the data and standard deviation, outliers present in data have a significant impact that reduces the spectrum of characteristic values.

Many machine learning algorithms may encounter issues due to these variations in the starting features. For algorithms that calculate distance, for instance, if any of the dataset's features have values having large or completely different ranges, that particular feature of the dataset will control the distance calculation.

The StandardScaler function of sklearn is based on the theory that the dataset's variables whose values lie in different ranges do not have an equal contribution to the model's fit parameters and training function and may even lead to bias in the predictions made with that model.

#Q.5 Consider statement “model = XGBClassifier(max_depth=5, objective='multi:softprob', n_estimators=1000, num_classes=10)” in the notebook explain purpose of each parameter of this constructor. What are we doing here defining a model with specific parameters or training the model?

#ANSWER:

#“model = XGBClassifier(max_depth=5, objective='multi:softprob', n_estimators=1000, num_classes=10)” This statement is used for model training and model fitting it determines and gives the parameter such as the maximum depth of the tree and no. of trees, categorize the class_id and also it gives the objective parameter essentially gives a fuzzy clustering for the distinct probability of belonging to each class.

#max_depth=5: This determines the maximum depth of the tree. In our case, we use a depth of two to make our decision tree. The default value is set to none. This will often result in over-fitted decision trees.

`#n_estimators=1000` `n_estimators`: The `n_estimators` parameter specifies the number of trees in the forest of the model. The default value for this parameter is 1000, which means that 10 different decision trees will be constructed in the random forest

`#objective='multi:softprob'`: The `multi:softprob` objective parameter essentially gives us a fuzzy clustering in which each observation is given a distinct probability of belonging to each class. In order to use these probabilities for classification, we will have to determine the max probability for each observation and assign a class.

`#num_classes=10` the `num_classes` still means the maximum `category_id` and this setting does not mean to include the background class. I checked the `dataset/create_pascal_tfrecord.py` and found the class label is the `category_id` (for pascal it ranges in `[1, 20]`). Setting `num_classes` to 21 merely adds an extra non-sense class, but does not influence the original 20 classes. The maximum `category_id` is 21 now, but there is no data of this class. Maybe this is a minor mistake. I suggest using the last setting `num_classes: 20`.

`#What step in ML pipeline fit function carries out?`

`#ANSWER:`

The `fit()` method takes the training data as arguments, which can be one array in the case of unsupervised learning, or two arrays in the case of supervised learning. Note that the model is fitted using `X` and `y`, but the object holds no reference to `X` and `y`. Pipelines allow us to streamline this process by compiling the preparation steps while easing the task of model tuning and monitoring. Scikit-Learn's Pipeline class provides a structure for applying a series of data transformations followed by an estimator. Step 1: Data Preprocessing. The first step in any pipeline is data preprocessing. Step 2: Data Cleaning. Next, this data flows to the cleaning step. Step 3: Feature Engineering. Step 4: Model Selection. Step 5: Prediction Generation. One definition of an ML pipeline is a means of automating the machine learning workflow by enabling data to be transformed and correlated into a model that can then be analyzed to achieve outputs. This type of ML pipeline makes the process of inputting data into the ML model fully automated. Before training a model, you should split your data into a training set and a test set. Each dataset will go through the data cleaning and preprocessing steps before you put it in a machine learning model.

It's not efficient to write repetitive code for the training set and the test set. This is when the scikit-learn pipeline comes into play.

Scikit-learn pipeline is an elegant way to create a machine learning model training workflow.